

# Dense and Sparse Matrices

*Dom Owens*

*12/11/2019*

## Dense Matrices

**Matrices** are 2-dimensional data structures. In R, entries are specified by column by default. Names can be reassigned without changing values.

```
x <- matrix(1:6, 2, 3) #2 by 3 matrix
x1 <- matrix(1:6, 3, 2, dimnames = list(c("X","Y","Z"), c("A","B"))) #give names to rows and columns
x
```

```
##      [,1] [,2] [,3]
## [1,]    1    3    5
## [2,]    2    4    6
```

```
x1
```

```
##    A B
## X 1 4
## Y 2 5
## Z 3 6
```

```
rownames(x1) <- c("R1","R2","R3") #add row names
x1
```

```
##    A B
## R1 1 4
## R2 2 5
## R3 3 6
```

We can preserve matrix dimensions when selecting rows with `drop=F`. Note that names are maintained.

```
x1[1,,drop=F] #select row 1
```

```
##    A B
## R1 1 4
```

**Arrays** are higher-dimensional data structures.

```
y <- array(1:8, c(2,2,2)) #3D data object
y #prints as transects
```

```
## , , 1
##
##      [,1] [,2]
## [1,]    1    3
## [2,]    2    4
##
## , , 2
##
##      [,1] [,2]
## [1,]    5    7
## [2,]    6    8
```

## Linear Systems

Matrix operations are useful for solving linear systems of the form  $Ax = b$ . Multiple approaches are available to do this.

We can try this on **Hilbert matrices**; these are close to singular, so are hard to invert.

`solve(A) %*% b` inverts  $A$  then multiplies by  $b$ .

Alternatively, we use `solve(A,b)`.

```
set.seed(123)
library(Matrix)
n <- 9
A <- as.matrix(Hilbert(n)) #generate Hilbert matrix
x <- matrix(runif(n), n, 1) #randomise x
b <- A%*%x # compute b
x1 <- solve(A) %*% b #method 1
x2 <- solve(A,b) #method 2
norm(x-x1, type = "1") # find numerical errors
```

```
## [1] 0.0001934015
```

```
norm(x - x2 ,type = "1")
```

```
## [1] 3.206041e-05
```

Clearly, the latter outperforms the former in terms of accuracy.

## Numerical Stability and Precision

Why has this happened? We should examine the computation underlying these divergences.

In R, floating point numbers are stored as *double precision* numbers. 64 bits store a representation of the number: the **sign** in 1 bit, the **exponent** (i.e. magnitude) in 11 bits, and **precision** in 52 bits. Hence have the largest number available:

```
2^1023 + 2^1022.999999999999999 #greatest number representable by R
```

```
## [1] 1.797693e+308
```

```
2^1024 #not representable
```

```
## [1] Inf
```

With small values, we can expect approximation errors.

```
1 + 1e-15 #not equal to 1
```

```
## [1] 1
```

```
print(1 + 1e-15, digits=22)
```

```
## [1] 1.0000000000000001110223
```

The operator `==`, which checks for equality, struggles for floating point digits.

```
0.1 + 0.2 - 0.3 == 0 #Should return TRUE
```

```
## [1] FALSE
```

From the above, we conclude that we should use `solve(A,b)` to minimise errors.

## Eigenvalues

`eigen` returns the **eigenvalues** and **eigenvectors** of the input matrix. We should specify the argument `symmetric = TRUE` if the input matrix is symmetric to reduce error in computation.

We demonstrate how the calculations differ below.

```
n <- 10
A1 <- matrix(rnorm(n*n), nrow=n, ncol=n) #n by n matrix of standard normal samples
A2 <- A1 + t(A1)
eA <- eigen(A2, symmetric=TRUE) #find eigenvalues
summary(abs(eA$values))

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.1392  1.2962  2.9755  3.5310  5.7445  8.2920

c(norm(eA$vectors %*% diag(eA$values) %*% t(eA$vectors) - A2, type='1'),
  norm(eA$vectors %*% t(eA$vectors) - diag(rep(1, n)), type='1')) #calculate divergence

## [1] 5.739853e-14 5.224826e-15
```

And with `symmetric = TRUE`

```
n <- 100
A1 <- matrix(rnorm(n*n), nrow=n, ncol=n); A2 <- A1 + t(A1)
eA <- eigen(A2, symmetric=TRUE) #find eigenvalues
summary(abs(eA$values))

##      Min.  1st Qu.   Median     Mean  3rd Qu.     Max.
## 0.008361  5.536487 11.357499 11.788913 17.657582 26.677891

c(norm(eA$vectors %*% diag(eA$values) %*% t(eA$vectors) - A2, type='1'),
  norm(eA$vectors %*% t(eA$vectors) - diag(rep(1, n)), type='1')) #calculate divergences

## [1] 2.705893e-12 1.954209e-13
```

These are larger by orders of magnitude. We see that specifying that the matrix is symmetric reduces precision error.

---

## Sparse Matrices

**Sparse matrices**, which consist mostly of zeros, are more of a challenge for a language to handle. It is not possible to store them in the same format as dense matrices due to memory constraints. `Matrix` stores dense matrices as `dgeMatrix` objects; `rankMatrix` will return the rank, and `rcond` returns the condition number; this quantifies the divergence between min and max eigenvalues.

```
rankMatrix(A) #A is full rank

## [1] 9
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
## attr(,"tol")
## [1] 1.998401e-15
```

```
rcond(A) #condition number
```

```
## [1] 9.093786e-13
```

Sparse matrices are stored by default as `dgCMatrix` objects.

We construct two sparse matrices and observe the memory difference:

```
nrows <- 1000
ncols <- 1000
vals <- sample(x=c(0, 1, 2), prob=c(0.98, 0.01, 0.01), size=nrows*ncols, replace=TRUE) #sample 1000*1000
m1 <- matrix(vals, nrow=nrows, ncol=ncols) # dense matrix representation
m2 <- Matrix(vals, nrow=nrows, ncol=ncols, sparse = TRUE) #sparse matrix representation
m1[1:2, 1:10]
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## [1,]    0    0    0    0    0    0    0    0    0    0
## [2,]    0    0    0    0    0    0    0    0    0    0
```

```
c(object.size(m1), object.size(m2)) #compare memory usage
```

```
## [1] 8000200 243448
```

The dense matrix `m1` uses 3 times as much space.

It is possible to coerce a `dgC` into a `dgT`, but not a `dgC` into a `dgR`

```
object.size(as(m2, 'dgMatrix')) # dgC` into a `dgT` possible
```

```
## 8001112 bytes
```

```
object.size(as(m1, 'dgRMatrix')) # `dgC` into a `dgR` not possible
```

```
## 243448 bytes
```

## Operations

Matrix operations on these objects may change or conserve the type.

Addition converts a `dcG` into a `dge`

```
B <- as(matrix(c(1,0,0,0,2,0), nrow=3, ncol=2), 'dgCMatrix')
B + 10
```

```
## 3 x 2 Matrix of class "dgeMatrix"
##      [,1] [,2]
## [1,]   11   10
## [2,]   10   12
## [3,]   10   10
```

Multiplication by a dense vector outputs a `dge` matrix

```
B %*% c(1,1)
```

```
## 3 x 1 Matrix of class "dgeMatrix"
##      [,1]
## [1,]    1
## [2,]    2
## [3,]    0
```

Multiplication by a sparse matrix, or taking a transpose, preserves sparsity.

```
B %*% Matrix(c(1, 1), nrow=2, ncol=1, sparse=TRUE)
```

```
## 3 x 1 sparse Matrix of class "dgCMatrix"
##
## [1,] 1
## [2,] 2
## [3,] .
```

```
t(B)
```

```
## 2 x 3 sparse Matrix of class "dgCMatrix"
##
## [1,] 1 . .
## [2,] . 2 .
```

## Solving large linear systems

Inverting sparse matrices is especially difficult, given the inverse of the matrix is not guaranteed to be sparse. We see this with a tridiagonal matrix

```
n <- 100
A1 <- bandSparse(n, k=c(0,1), diag=list(rep(1,n), rep(-0.2, n)), symm=T) #Tridiagonal
A1inv <- solve(A1)
c(object.size(A1), object.size(A1inv))
```

```
## [1] 11488 121824
```

```
sum(abs(A1inv) < 1e-100) #no entries near to 0
```

```
## [1] 0
```